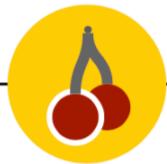


# **CherryPick**: Adaptively Unearthing the Best Cloud Configurations for Big Data Analytics

**Omid Alipourfard**, Hongqiang Harry Liu, Jianshu Chen,  
Shivaram Venkataraman, Minlan Yu, Ming Zhang



**Yale**

Microsoft®  
**Research**

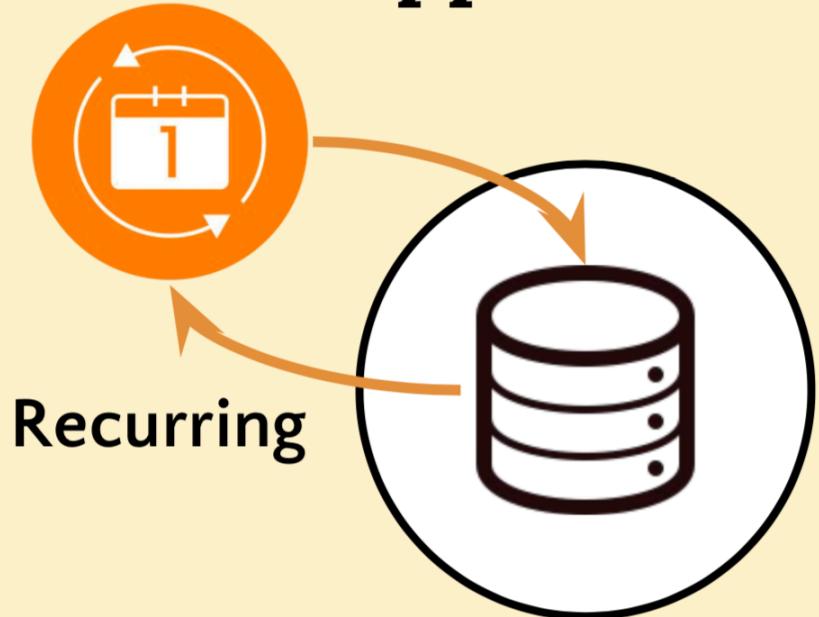
**Berkeley**  
UNIVERSITY OF CALIFORNIA



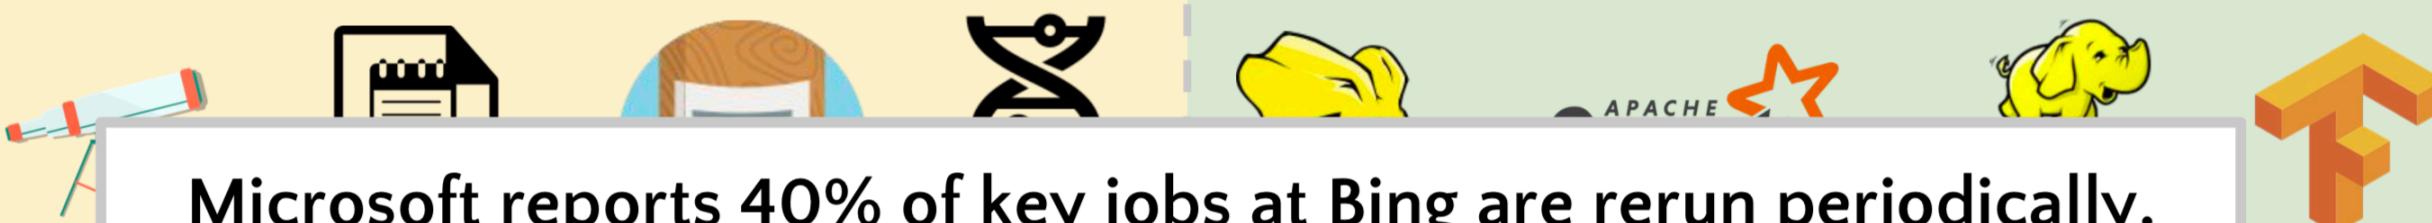
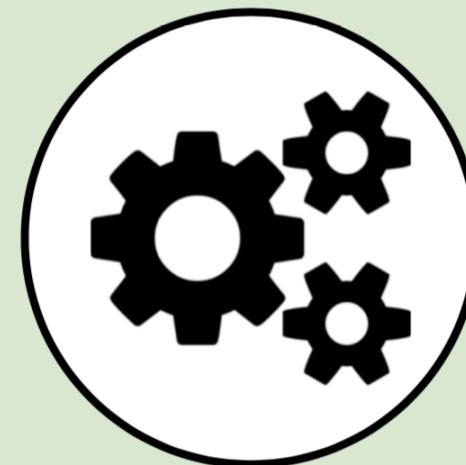
The logo for Alibaba.com, featuring the company name in orange lowercase letters next to a stylized orange 'e' symbol.

# Recurring jobs are popular

## Apps & Data



## Frameworks



Microsoft reports 40% of key jobs at Bing are rerun periodically.

## Providers



## Machine Type

r3.8xlarge, i2.8xlarge,  
m4.8xlarge, c4.8xlarge,

## Cluster Size

Hundreds of instance types and instance count combinations

## Azure



A12, D1, D2, D3, L4s, ...

n1-standard-4, n1-highmem-2,  
n1-highcpu-4, f1-micro, ...  
+ configurable VMs

# Choosing a good configuration is important

## Better performance:

- For the same cost: best/worst running time is up to 3x
  - Worst case has good CPUs whereas the memory is bottlenecked.

## Lower cost:

- For the same performance: best/worst cost is up to 12x
  - No need for expensive dedicated disks in the worst config.
  - 2\$ vs. 24\$ per job with 100s of monthly runs

Application	Avg/min	Max/min
TPC-DS	3.4	9.6
TPC-H	2.9	12
Regression (SparkML)	2.6	5.2
TeraSort	1.6	3.0

**How to find the best cloud configuration**

One that minimizes the cost given a performance constraint

**for a recurring job, given its representative workload?**

## High Accuracy

Close the optimal configurations

## Adaptivity

Works across all big-data apps

## Key Challenges

Modeling

Searching

## Low Overhead

Only runs a few configuration

# Existing solution: searching

- Systematically search each dimension (Coordinate descent)
  - On each resources: RAM, CPU, disk, cluster sizes
- Problem: not accurate
  - Non-convex performance/cost curves across many resources
  - If you search one dimension, drop early, it would mislead you later

# Existing solution: modeling

- Modeling the resource-perf/cost tradeoffs
  - Ernest [NSDI 16] models machine learning apps for each machine type
- Problem: **not adaptive**
  - Frameworks (e.g., MapReduce or Spark)
  - Applications (e.g., machine learning or database)
  - Machine type (e.g., memory vs CPU intensive)

**High Accuracy**

**Exhaustive**

High overhead

**Adaptivity**

**CherryPick**

**Modeling**

Ernest

Not general

**Searching**

Coord. Descent

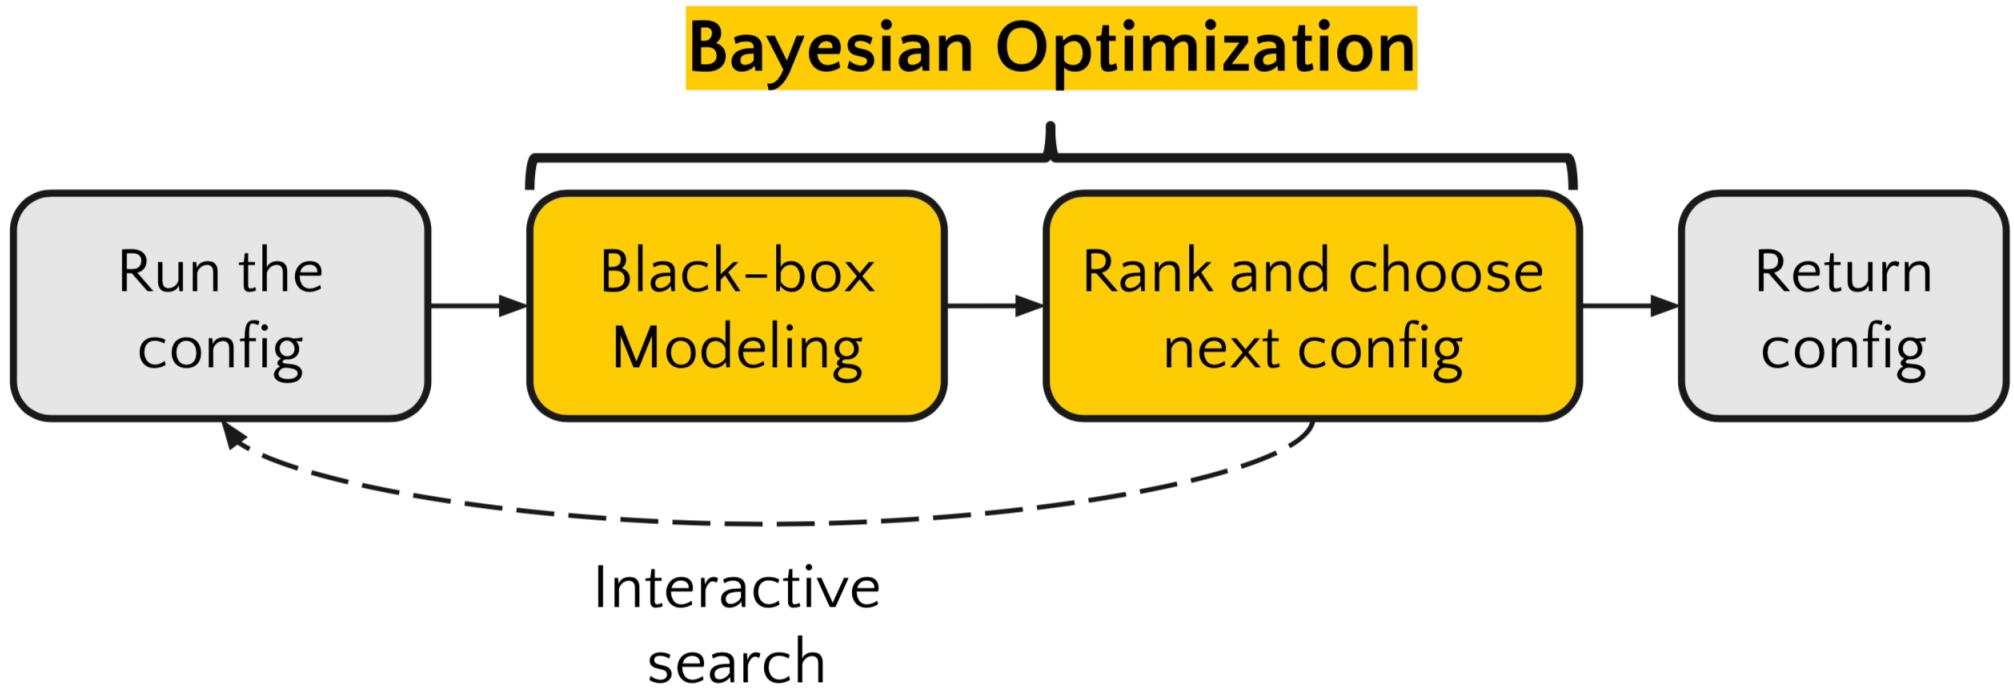
Not accurate

**Low Overhead**

# Key idea of CherryPick

- **Adaptivity: black-box modeling**
  - Without knowing the structure of each application
- **Accuracy: modeling for ranking configurations**
  - No need to be accurate everywhere
- **Low overhead: interactive searching**
  - Smartly select next run based on existing runs

# Workflow of CherryPick

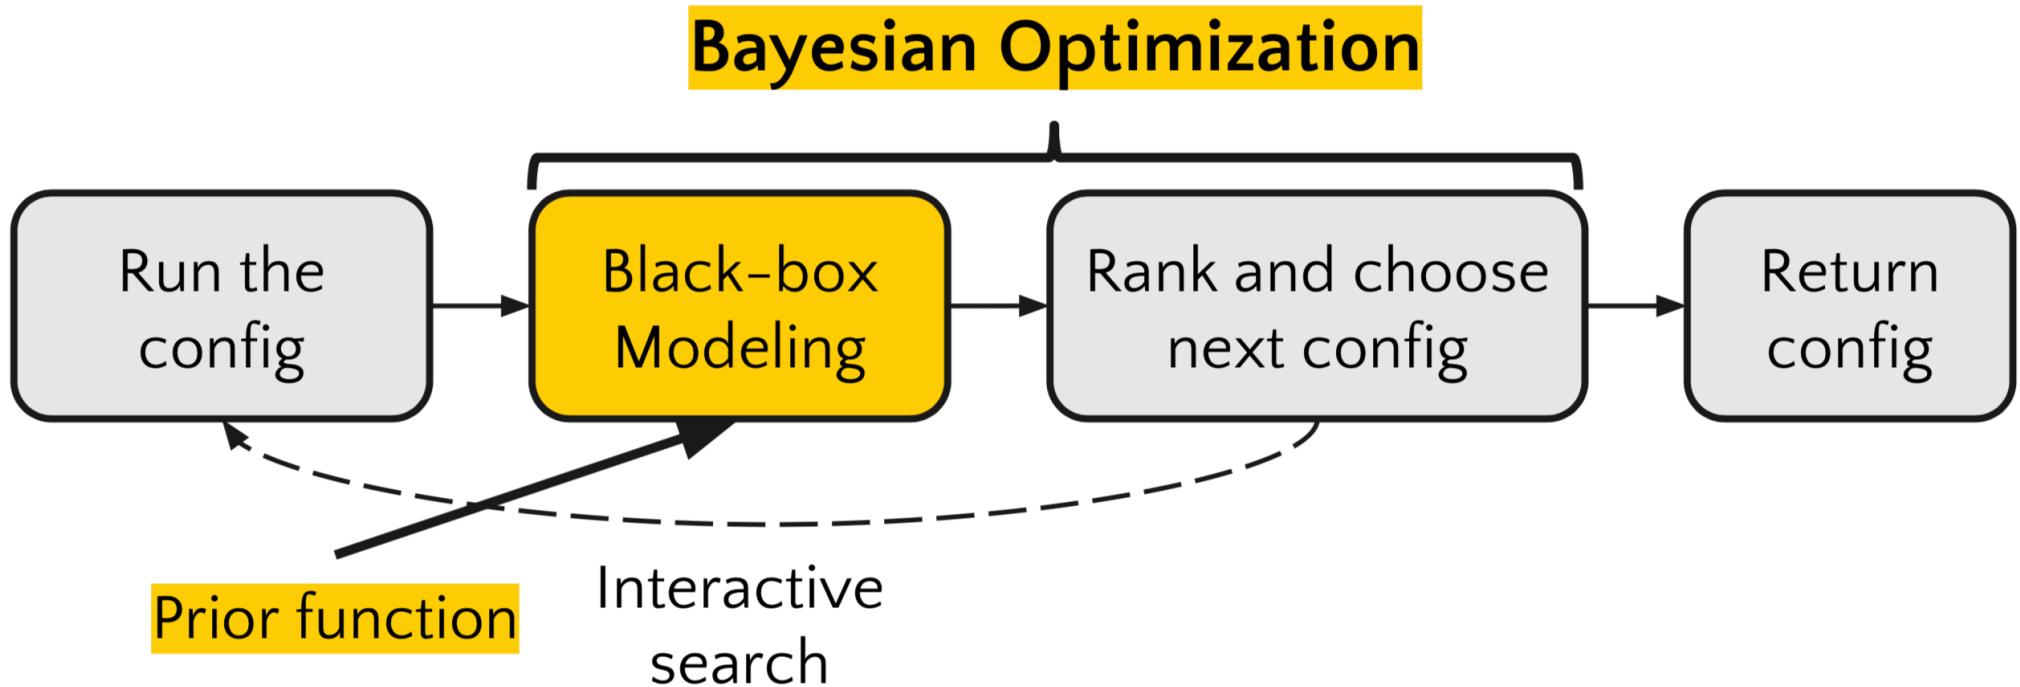


# Workflow of CherryPick

$$\underset{\vec{x}}{\text{minimize}} \quad C(\vec{x}) = P(\vec{x}) \times T(\vec{x})$$

$$\text{subject to} \quad T(\vec{x}) \leq \mathcal{T}_{max}$$

# Workflow of CherryPick



# Workflow of CherryPick

- Noises are common in the cloud
  - Shared environment means inherent noise from other tenants
  - Even more noisy under failures, stragglers
  - **Strawman solution:** run multiple times, but high overhead.
  -
- Bayesian optimization is good at handling additive noise
  - Merge the noise in the confidence interval

$$f(\vec{X}) + \epsilon \leftarrow \text{Noise}$$

(learned by monitoring or historical data)

# Workflow of CherryPick

A lot of the noise in cloud is multiplicative

- Example: if VMs IO slows down due to overloading
  - Writing 1G to disk (5 sec normally) now takes 6 secs (by 20%)
  - Writing 10G to disk (50 sec normally) now takes 60 secs (by 20%)

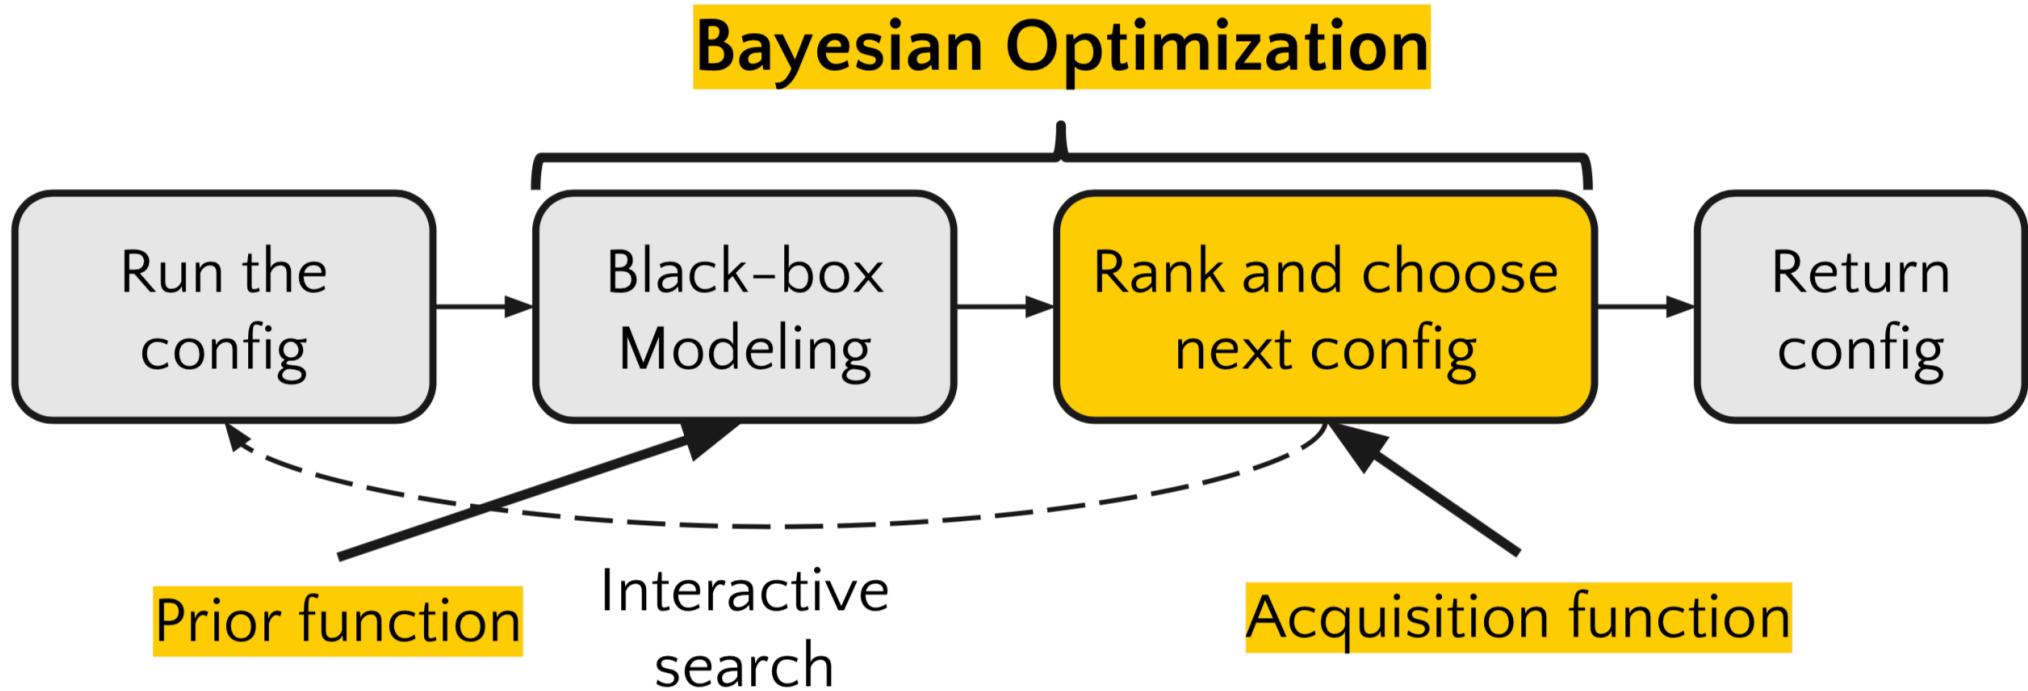
$$\hat{C}(\vec{X}) = C(\vec{X}) \times (1 + \epsilon)$$

**The max. relative variance in a given cloud.**

Use the log function to change to additive noise

$$\log \hat{C}(\vec{X}) = \log C(\vec{X}) + \log(1 + \epsilon)$$

# Workflow of CherryPick



# Workflow of CherryPick

$$EI(\vec{x})' = P[T(\vec{x}) \leq \mathcal{T}_{max}] \times EI(\vec{x})$$

# Further customizations

- **Discretize features:**
  - Deal with infeasible configs and large searching space
  - Discretize the feature space
- **Stopping condition:**
  - Trade-off between accuracy and searching cost
  - Use the acquisition function knob
- **Starting condition:**
  - Should fully cover the whole space
  - Use quasi-random search

# Evaluation

## 5 big-data benchmarks

- Database:
  - TPC-DS
  - TPC-H
- MapReduce:
  - TeraSort
- Machine Learning:
  - SparkML Regression
  - SparkML Kmeans

## 66 cloud configurations

- 30 GB-854 GB RAM
- 12-112 cores
- 5 machine types

# Evaluation

## Metrics

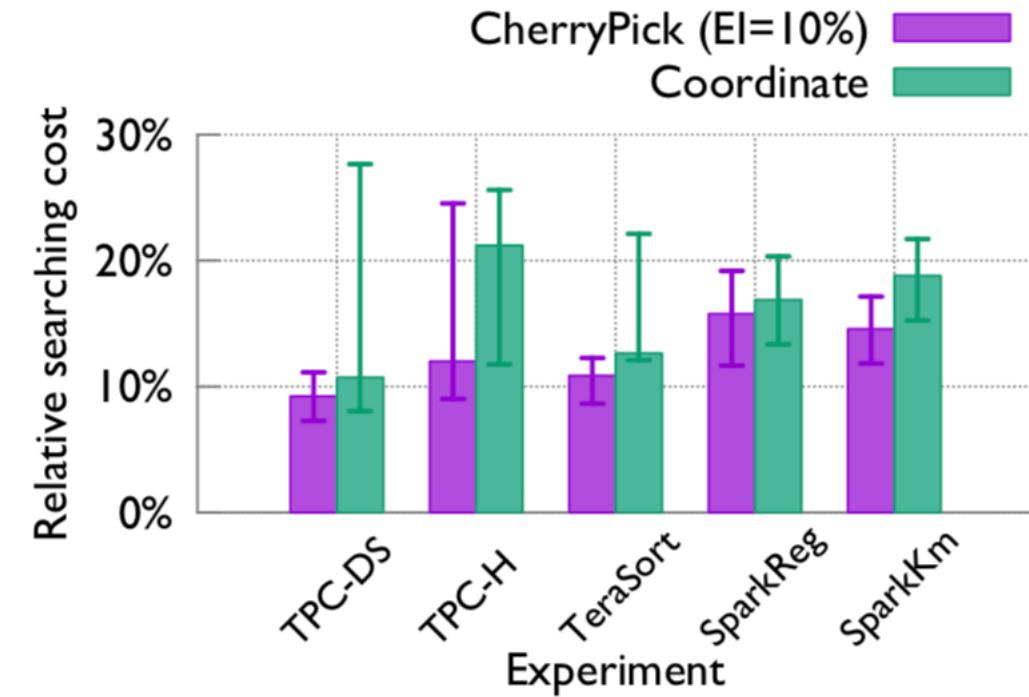
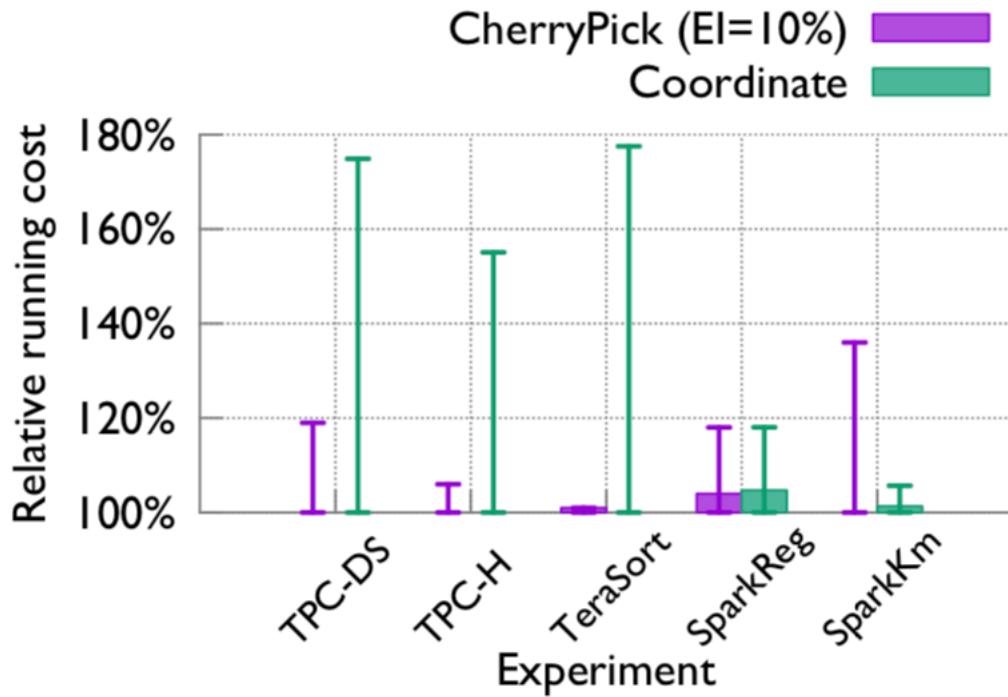
- Accuracy: running cost compared to the optimal configuration
- Overhead: searching cost of suggesting a configuration

## Comparing with

- Searching: random search, coordinate descent
- Modeling: Ernest [NSDI'16]

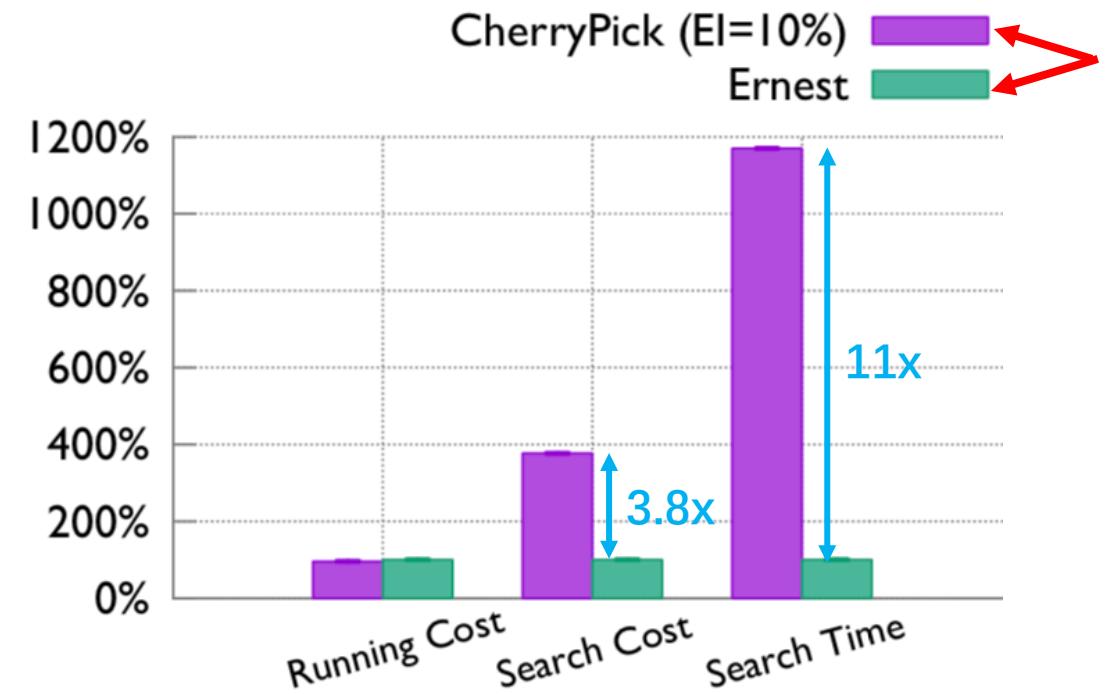
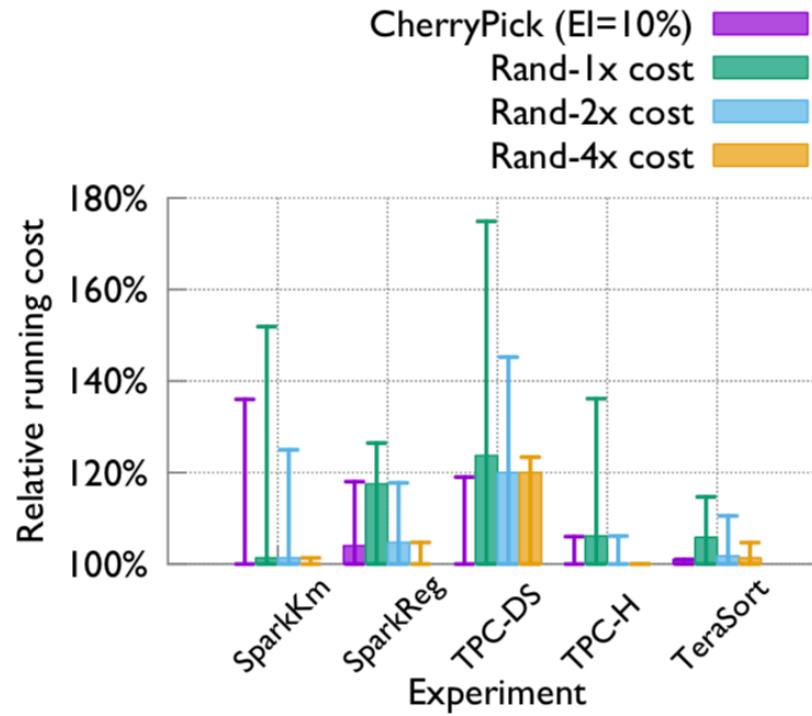
# Evaluation

- CD: Finds opt/near-opt conf with low search cost and time
- CD: higher chance, more stable



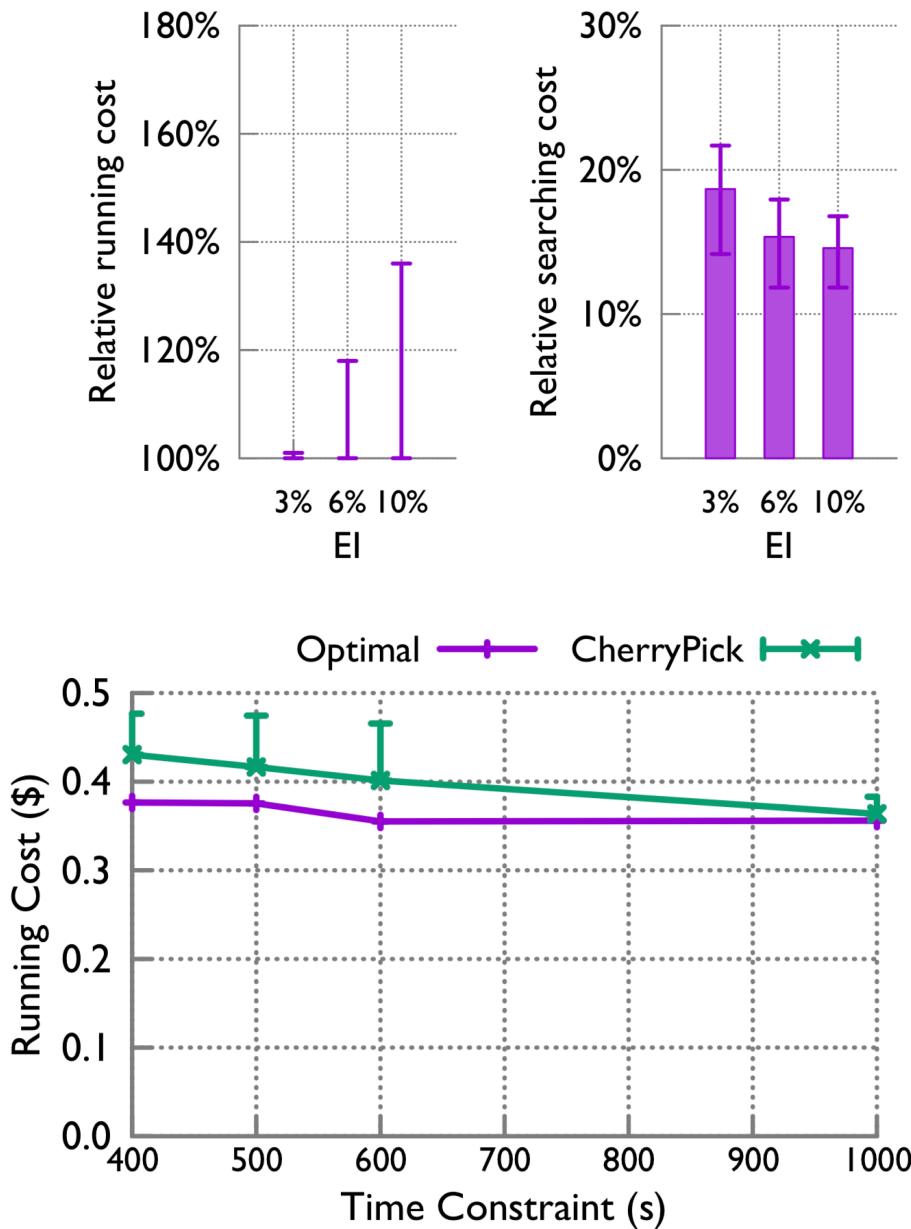
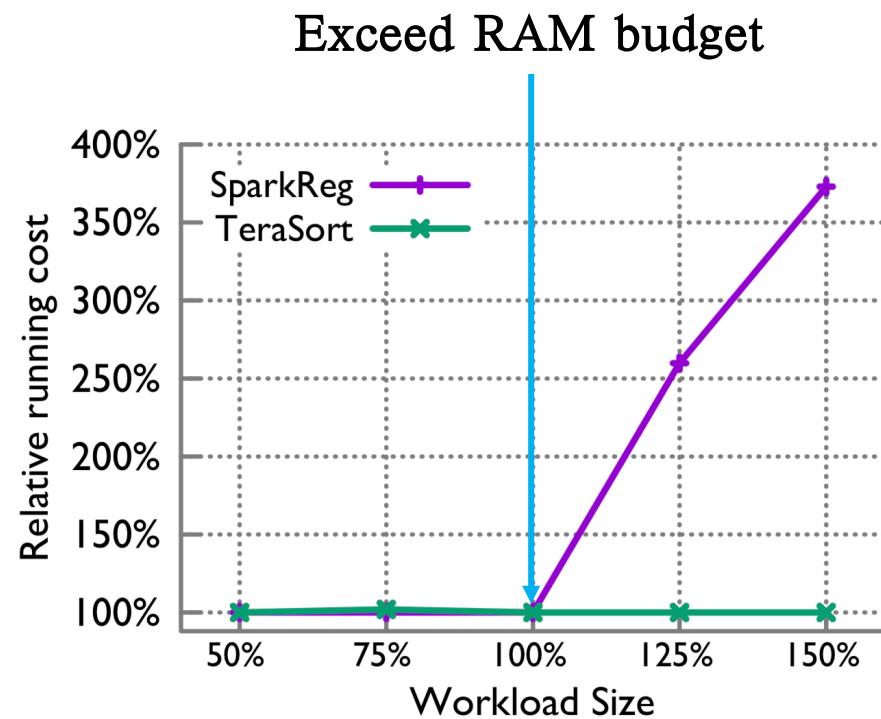
# Evaluation

- Random: better conf & more stability with similar budget
- Ernest: similar running cost, lower search cost and time



# Evaluation

- Stable against EI and time constraint
- Rethink representativeness of workload



# Conclusion

Adaptivity

Low overhead

High accuracy

**Black-box modeling:**

Requires running the cloud configuration.

**Restricted amount of information:**

Only a few runs available.

“... do not solve a more general problem ...  
try to get the answer that you really need”  
—Vladimir Vapnik