

Neural Architecture Optimization

Shen Yu 1901111296

2019.7.11

Motivation

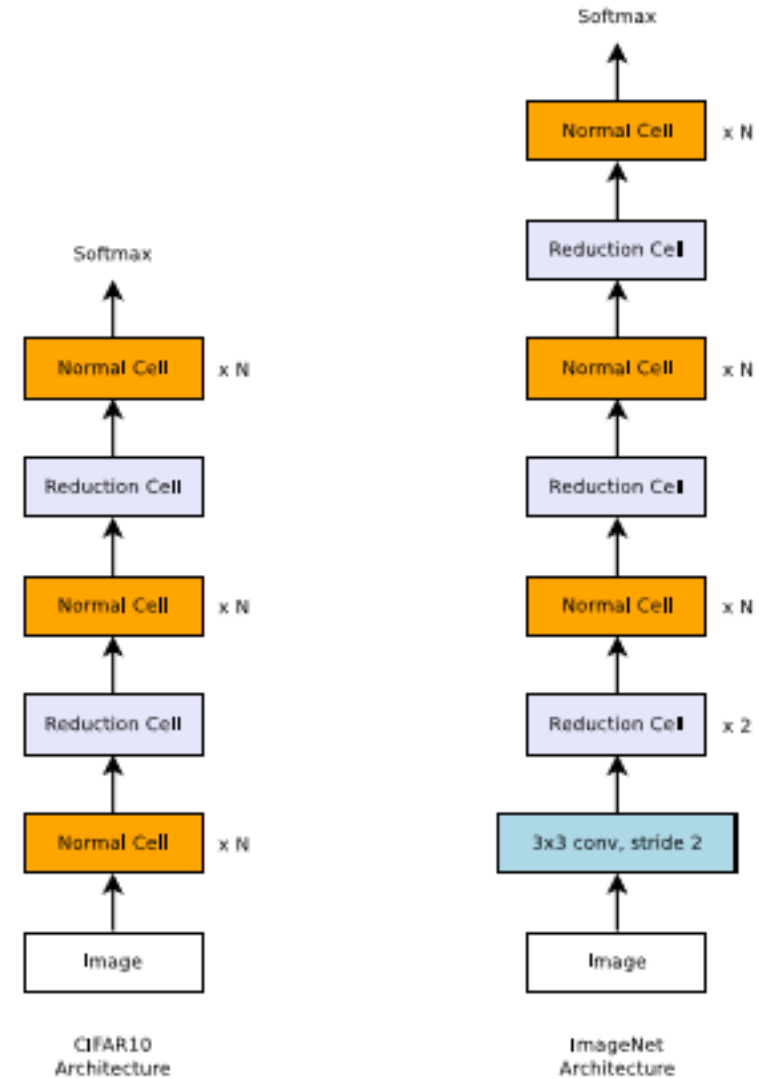
- Substantial effort of human experts on discovering architectures
- Searching the best architecture within discrete space is inefficient

Contributions

- Propose to optimize network architecture by mapping architectures into a continuous vector space
- Achieve improved efficiency in discovering powerful convolutional and recurrent architectures

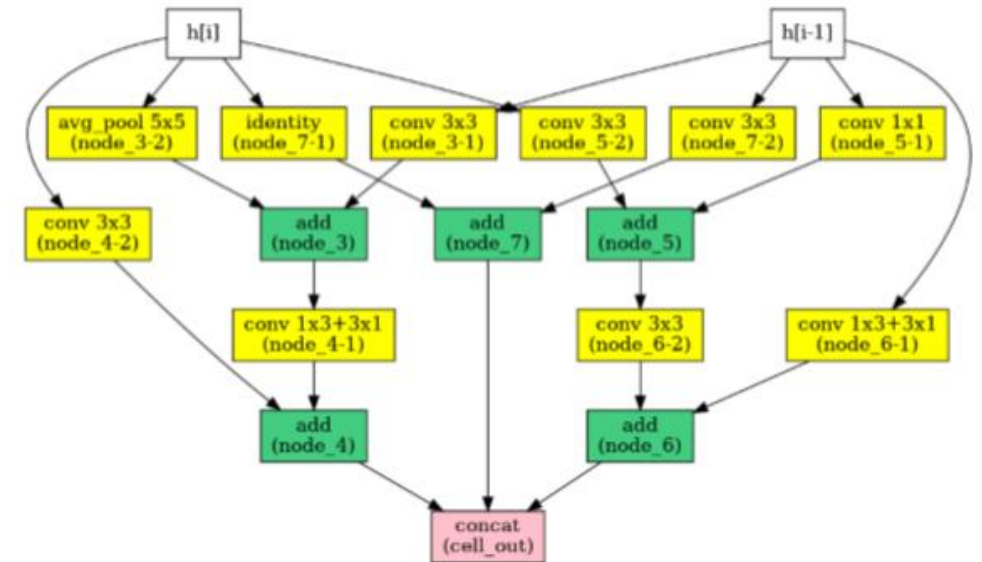
Related Work

- Reinforcement learning
 - ENAS
- Evolutionary algorithm
 - AmoebaNet
- SMBO
 - PNAS



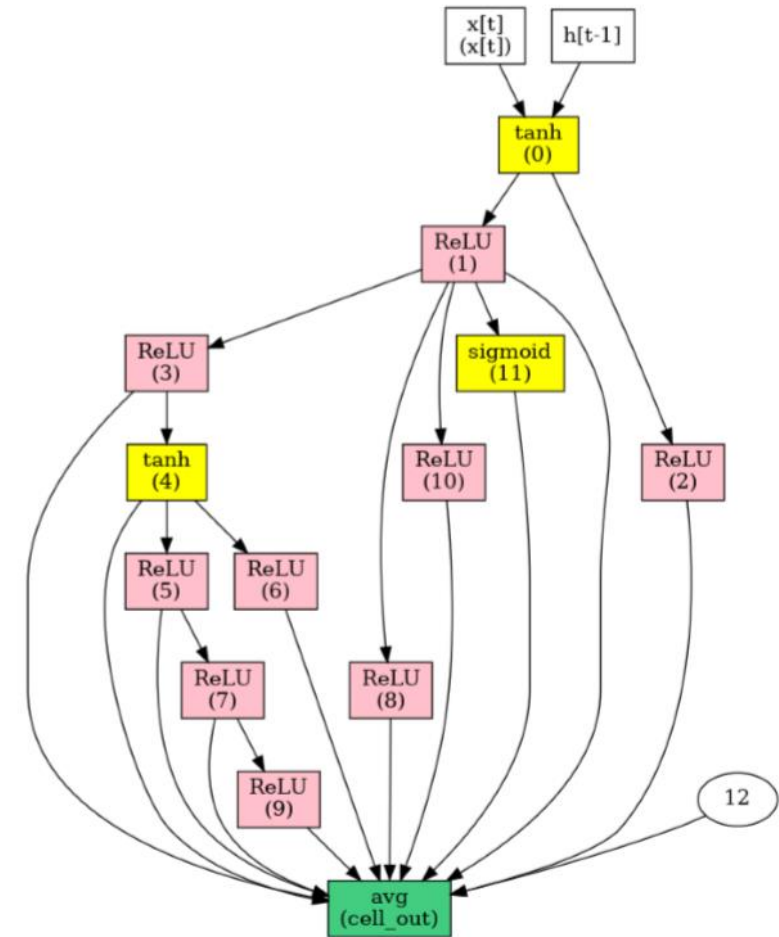
Architecture space (CNN)

- 2 cells (Normal and Reduction)
 - B(=5) blocks
 - Input 1 from two previous cells or previous blocks
 - Input 2 from two previous cells or previous blocks
 - Operation applied to input 1
 - Operation applied to input 2
- Output the concatenation of outputs of unused blocks




Architecture space (RNN)

- An RNN cell
 - B(=12) nodes
 - Input from previous nodes
 - Activation
 - Output the average of the outputs of all the nodes



Overview



Encoder

- Input: A sequence $\{x_1 x_2 \dots x_T\}$
 - $x_{example} = \text{"node1 conv 3x3 node2 max-pooling 3x3"}$
- Output: Architecture embedding e_x
 - $\{h_1 h_2 \dots h_T\} \in R^{T \times d}$, h_t is the hidden state at t-th timestep
- Architecture: A single layer LSTM with d hidden units

Performance Predictor

- Input: Mean pooling of embedding e_x

$$\bar{e}_x = \frac{1}{T} \sum_{t=1}^T h_t$$

- Output: Performance prediction
- Architecture: FFNN

Decoder

- Input: Architecture embedding e_x
- Output: Predicted architecture x'
- Architecture: LSTM with attention

Loss function

- Performance predictor

$$L_{pp} = (s_x - f(E(x)))^2$$

- Decoder

$$L_{rec} = -\log P(x|E(x)) = -\sum_{t=1}^T \log P(x_t|E(x), x_{<t}) = -\sum_{t=1}^T \log \frac{\exp(W_{x_t})}{\sum_{x' \in V_t} \exp(W_{x'})}$$

- Final loss

$$L_{final} = \lambda L_{pp} + (1 - \lambda) L_{rec}$$

Algorithm

Algorithm 1 Neural Architecture Optimization

Input: Initial candidate architectures set X to train NAO model. Initial architectures set to be evaluated denoted as $X_{eval} = X$. Performances of architectures $S = \emptyset$. Number of seed architectures K . Step size η . Number of optimization iterations L .

for $l = 1, \dots, L$ **do**

Train each architecture $x \in X_{eval}$ and evaluate it to obtain the dev set performances $S_{eval} = \{s_x\}, \forall x \in X_{eval}$. Enlarge S : $S = S \cup S_{eval}$.

Train encoder E , performance predictor f and decoder D by minimizing Eqn.(1), using X and S .

Pick K architectures with top K performances among X , forming the set of seed architectures X_{seed} .

For $x \in X_{seed}$, obtain a better representation $e_{x'}$ from e_x using Eqn. (2), based on encoder E and performance predictor f . Denote the set of enhanced representations as $E' = \{e_{x'}\}$.

Decode each x' from $e_{x'}$ using decoder, set X_{eval} as the set of new architectures decoded out: $X_{eval} = \{D(e_{x'}), \forall e_{x'} \in E'\}$. Enlarge X as $X = X \cup X_{eval}$.

end for

Output: The architecture within X with the best performance

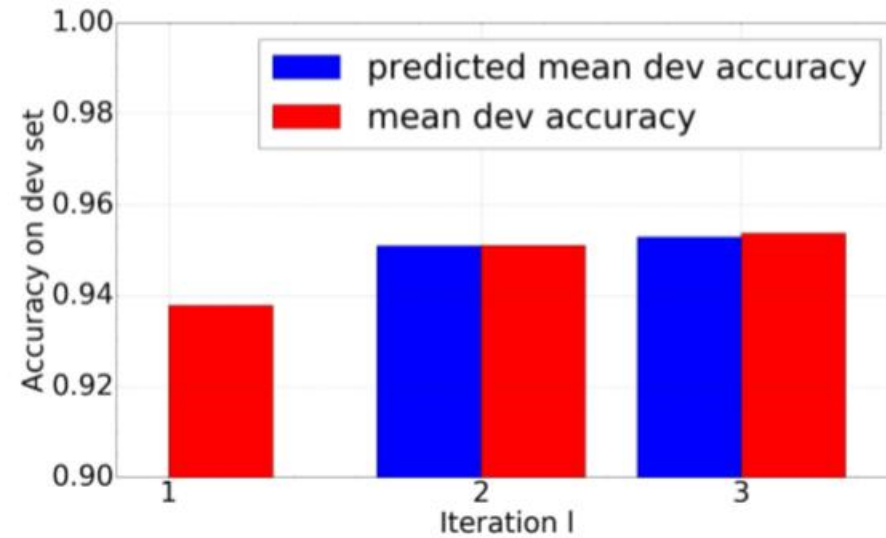
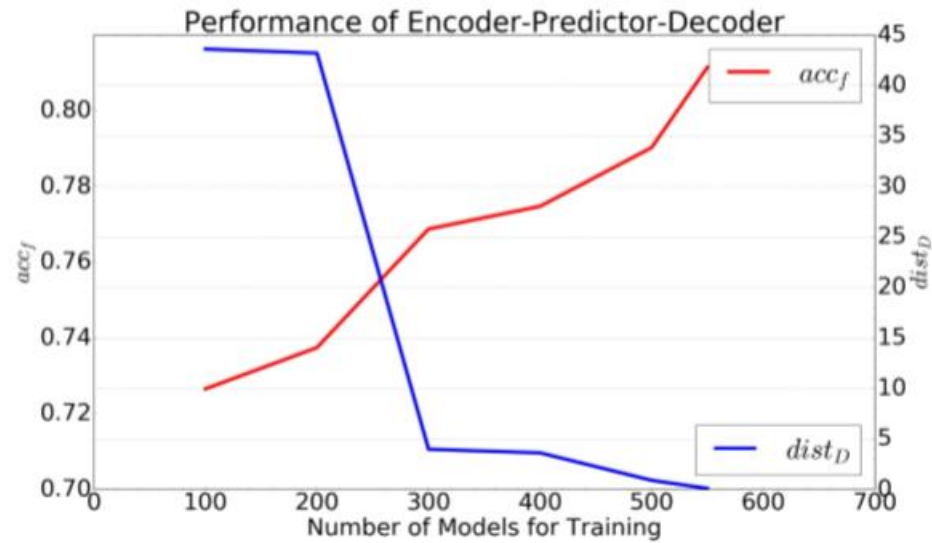
Trick

- Data augmentation
 - For each (x_1, s_x) , add an additional pair (x_2, s_x) where x_2 is symmetrical to x_1 and use both pairs to train the encoder and performance predictor
- $x_{example1}$ = “node1 conv 3x3 node2 max-pooling 3x3”
- $x_{example2}$ = “node2 max-pooling 3x3 node1 conv 3x3”
- $s(x_{example1}) = s(x_{example2})$

Performance on Cifar-10

Model	B	N	F	#op	Error(%)	#params	M	GPU Days
DenseNet-BC [19]		100	40	/	3.46	25.6M	/	/
ResNeXt-29 [43]				/	3.58	68.1M	/	/
NASNet-A [47]	5	6	32	13	3.41	3.3M	20000	2000
NASNet-B [47]	5	4	N/A	13	3.73	2.6M	20000	2000
NASNet-C [47]	5	4	N/A	13	3.59	3.1M	20000	2000
Hier-EA [27]	5	2	64	6	3.75	15.7M	7000	300
AmoebaNet-A [38]	5	6	36	10	3.34	3.2M	20000	3150
AmoebaNet-B [38]	5	6	36	19	3.37	2.8M	27000	3150
AmoebaNet-B [38]	5	6	80	19	3.04	13.7M	27000	3150
AmoebaNet-B [38]	5	6	128	19	2.98	34.9M	27000	3150
AmoebaNet-B + Cutout [38]	5	6	128	19	2.13	34.9M	27000	3150
PNAS [26]	5	3	48	8	3.41	3.2M	1280	225
ENAS [36]	5	5	36	5	3.54	4.6M	/	0.45
Random-WS	5	5	36	5	3.92	3.9M	/	0.25
DARTS + Cutout [28]	5	6	36	7	2.83	4.6M	/	4
NAONet	5	6	36	11	3.18	10.6M	1000	200
NAONet	5	6	64	11	2.98	28.6M	1000	200
NAONet + Cutout	5	6	128	11	2.11	128M	1000	200
NAONet-WS	5	5	36	5	3.53	2.5M	/	0.3

Analysis on NAO



Performance on Cifar-100

Model	B	N	F	#op	Error (%)	#params
DenseNet-BC [19]	/	100	40	/	17.18	25.6M
Shake-shake [15]	/	/	/	/	15.85	34.4M
Shake-shake + Cutout [11]	/	/	/	/	15.20	34.4M
NASNet-A [47]	5	6	32	13	19.70	3.3M
NASNet-A [47] + Cutout	5	6	32	13	16.58	3.3M
NASNet-A [47] + Cutout	5	6	128	13	16.03	50.9M
PNAS [26]	5	3	48	8	19.53	3.2M
PNAS [26] + Cutout	5	3	48	8	17.63	3.2M
PNAS [26] + Cutout	5	6	128	8	16.70	53.0M
ENAS [36]	5	5	36	5	19.43	4.6M
ENAS [36] + Cutout	5	5	36	5	17.27	4.6M
ENAS [36] + Cutout	5	5	36	5	16.44	52.7M
AmoebaNet-B [38]	5	6	128	19	17.66	34.9M
AmoebaNet-B [38] + Cutout	5	6	128	19	15.80	34.9M
NAONet + Cutout	5	6	36	11	15.67	10.8M
NAONet + Cutout	5	6	128	11	14.75	128M

Performance on PTB

Models and Techniques	#params	Test Perplexity	GPU Days
Vanilla LSTM [45]	66M	78.4	/
LSTM + Zoneout [23]	66M	77.4	/
Variational LSTM [14]	19M	73.4	
Pointer Sentinel-LSTM [33]	51M	70.9	/
Variational LSTM + weight tying [20]	51M	68.5	/
Variational Recurrent Highway Network + weight tying [46]	23M	65.4	/
4-layer LSTM + skip connection + averaged weight drop + weight penalty + weight tying [31]	24M	58.3	/
LSTM + averaged weight drop + Mixture of Softmax + weight penalty + weight tying [44]	22M	56.0	/
NAS + weight tying [47]	54M	62.4	1e4 CPU days
ENAS + weight tying + weight penalty [36]	24M	58.6 ⁵	0.5
Random-WS + weight tying + weight penalty	27M	58.81	0.4
DARTS+ weight tying + weight penalty [28]	23M	56.1	1
NAONet + weight tying + weight penalty	27M	56.0	300
NAONet-WS + weight tying + weight penalty	27M	56.6	0.4

Q&A